

DL ASSIGNMENT 3 PART1

[\[DRIVE LINK FOR SAVED MODELS\]](#), [\[COLAB LINK FOR THE CODE FILE\]](#)

A. Implement a CNN architecture with

- block1 followed by FC layers, and a softmax layer
- block1, 2 followed by FCs, and a softmax layer
- block 1,2, 3, followed by FCs and a softmax layer

```
CNN(  
  (block1): Sequential(  
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (4): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (block2): Sequential(  
    (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (4): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (block3): Sequential(  
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (4): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Linear(in_features=1024, out_features=512, bias=True)  
    (1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=512, out_features=128, bias=True)  
    (4): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=128, out_features=10, bias=True)  
  )  
)
```

- For part a, consider block1 only
- For part b, consider block1 and block2 only.
- For part c, consider the whole architecture.

B. For all the three architectures apply the Tanh or ReLU activation function on all layers.

- ReLU was chosen.

C. Implement Dropout and use

- After convolutional layers
Dropouts were tried after every block of convolution layers, but Batch Normalization gave better results.
- Between FC layers
Dropouts have been applied after all the fully connected layers except the output layer.

DL ASSIGNMENT 3 PART1

1. Visualize 10 random images from each class.

- Dolphin

Class 0:



- Flower

Class 1:




- Keyboard

Class 2:




- Tree

Class 3:




- Building

Class 4:



- House

Class 5:



DL ASSIGNMENT 3 PART1

- Forest

Class 6:



- Bike

Class 7:



- Tree

Class 8:



- Fishing

Class 9:



2. Analyze the accuracy and loss while adding block 1, block 2, and block 3 with mentioned non-linearities.

- Using just block1 gives a not very bad model.
- Using block1 and block2, the model starts overfitting so we added BatchNorm and L2 norms.
- Using all 3 blocks, we got an improved model hence we selected this architecture.

DL ASSIGNMENT 3 PART1

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/block_1_only.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Tra:
```

Epochs=84, Train-Loss: 0.3532, Val-Loss: 0.5595, Train-Acc: 87.44 %, Val-Acc: 81.07 %

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/block_1_and_2.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Tra
```

Epochs=76, Train-Loss: 0.2198, Val-Loss: 0.5322, Train-Acc: 92.49 %, Val-Acc: 82.00 %

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/block_1_2_and_3.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Trai
```

Epochs=60, Train-Loss: 0.2010, Val-Loss: 0.5346, Train-Acc: 92.89 %, Val-Acc: 83.07 %

DL ASSIGNMENT 3 PART1

3. Analyze the accuracy and loss while changing the dropout probability. Report the results obtained on CNN with all three blocks. Try at least 3 different dropout probabilities.

- All the dropout probabilities were set to 0.15, 0.25, and 0.50. Results have been shown below.
- We used the 3 block architecture that was found to be the best from the above analysis.
- From the following analysis, we found that **dropout=0.5** gave the best results and this will be our **best model** for the subsequent analysis.

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/dropout_0_15.pt"
```

```
checkpoint = torch.load(PATH)
```

```
model = CNN().cuda()
```

```
model.load_state_dict(checkpoint['model_state_dict'])
```

```
epoch = checkpoint['epoch']
```

```
train_loss = checkpoint['train_loss']
```

```
val_loss = checkpoint['val_loss']
```

```
train_acc = checkpoint['train_acc']
```

```
val_acc = checkpoint['val_acc']
```

```
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Tra
```

```
Epochs=42, Train-Loss: 0.1866, Val-Loss: 0.5679, Train-Acc: 93.11 %, Val-Acc: 81.20 %
```

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/dropout_0_25.pt"
```

```
checkpoint = torch.load(PATH)
```

```
model = CNN().cuda()
```

```
model.load_state_dict(checkpoint['model_state_dict'])
```

```
epoch = checkpoint['epoch']
```

```
train_loss = checkpoint['train_loss']
```

```
val_loss = checkpoint['val_loss']
```

```
train_acc = checkpoint['train_acc']
```

```
val_acc = checkpoint['val_acc']
```

```
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Tra
```

```
Epochs=53, Train-Loss: 0.1560, Val-Loss: 0.5363, Train-Acc: 94.59 %, Val-Acc: 82.40 %
```

DL ASSIGNMENT 3 PART1

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/dropout_0_50.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Trai

Epochs=57, Train-Loss: 0.2125, Val-Loss: 0.4951, Train-Acc: 92.59 %, Val-Acc: 84.13 %
```

4. Report the best accuracy with model architecture and detailed analysis of choosing specific hyperparameters, different training techniques, and data augmentation used (if any).

The best model: test accuracy = [84.13%](#), training accuracy = [92.59%](#).

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/best_model.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Trai

Epochs=57, Train-Loss: 0.2125, Val-Loss: 0.4951, Train-Acc: 92.59 %, Val-Acc: 84.13 %
```

[The Best Model's Architecture:](#)

DL ASSIGNMENT 3 PART1

```
CNN(  
  (block1): Sequential(  
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (block2): Sequential(  
    (0): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (block3): Sequential(  
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Linear(in_features=1024, out_features=512, bias=True)  
    (1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=512, out_features=128, bias=True)  
    (5): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): ReLU(inplace=True)  
    (7): Dropout(p=0.5, inplace=False)  
    (8): Linear(in_features=128, out_features=10, bias=True)  
  )  
)
```

Hyperparameters:

- Batch-Size=64
- Epochs=100
- L2 Regularization, $\lambda = 1e-5$

Training Techniques:

- Adam Optimizer with default learning rate=0.001 and L2 regularization of $1e-5$.

```
optimizer = torch.optim.Adam(model.parameters(), weight_decay=0.00001)
```

Data Augmentation:


DL ASSIGNMENT 3 PART1

```
train_transform = T.Compose([
    T.ToPILImage(),
    T.RandomHorizontalFlip(),
    T.RandomCrop(32, padding = 4),
    T.ToTensor(),
])
```

- RandomHorizontalFlip()
 - creating additional samples by horizontally flipping original samples.
- RandomCrop(32, padding = 4)
 - creating additional samples by first padding the original samples and then taking a random crop of the same size as the original image sample.


5. Visualize some convolutional filters and feature maps obtained after each block.

Input Image
Class: 4




Feature Maps

- After Block1




- After Block2




- After Block3

DL ASSIGNMENT 3 PART1




Convolution Filters

- Block1 - convX1




- Block1 - convX2



- Block2 - convY1



- Block2 - convY2



- Block3 - convZ1



- Block3 - convZ2



6. Analyze the results of the best model when all the activation functions are removed. Justify the performance drop.

- Train accuracy dropped from 92 to 84 while test accuracy dropped from 84 to 80.
- When we drop the activation functions, indirectly we are switching to the linear activation function.
- We already know that linear activation functions render the hidden layers useless since the hypothesis generated can be expressed as just a linear function of the inputs as we saw in one of the questions asked in the mid-sem examination.
- Without Activations Model Evaluations

DL ASSIGNMENT 3 PART1

```
PATH = "/content/drive/MyDrive/DL_A3/A3_P1_Models/best_model_no_acti.pt"

checkpoint = torch.load(PATH)

model = CNN().cuda()
model.load_state_dict(checkpoint['model_state_dict'])

epoch = checkpoint['epoch']
train_loss = checkpoint['train_loss']
val_loss = checkpoint['val_loss']
train_acc = checkpoint['train_acc']
val_acc = checkpoint['val_acc']
print (f'Epochs={epoch+1}, Train-Loss: {train_loss:.4f}, Val-Loss: {val_loss:.4f}, Trai

Epochs=53, Train-Loss: 0.4378, Val-Loss: 0.5902, Train-Acc: 84.85 %, Val-Acc: 80.53 %
```

7. [Bonus] During the demo you are given labels of test data (format will be the same as training data), you have to evaluate the test accuracy of your best model.

<https://colab.research.google.com/drive/1r4TOgR0X2woRziqrib6hGBZ6AE9jhosu?usp=sharing>

References:

- <https://www.youtube.com/playlist?list=PLqnsIRFeH2UrcDBWF5mfPGpqQDSta6VK4>
- <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>
- <https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>